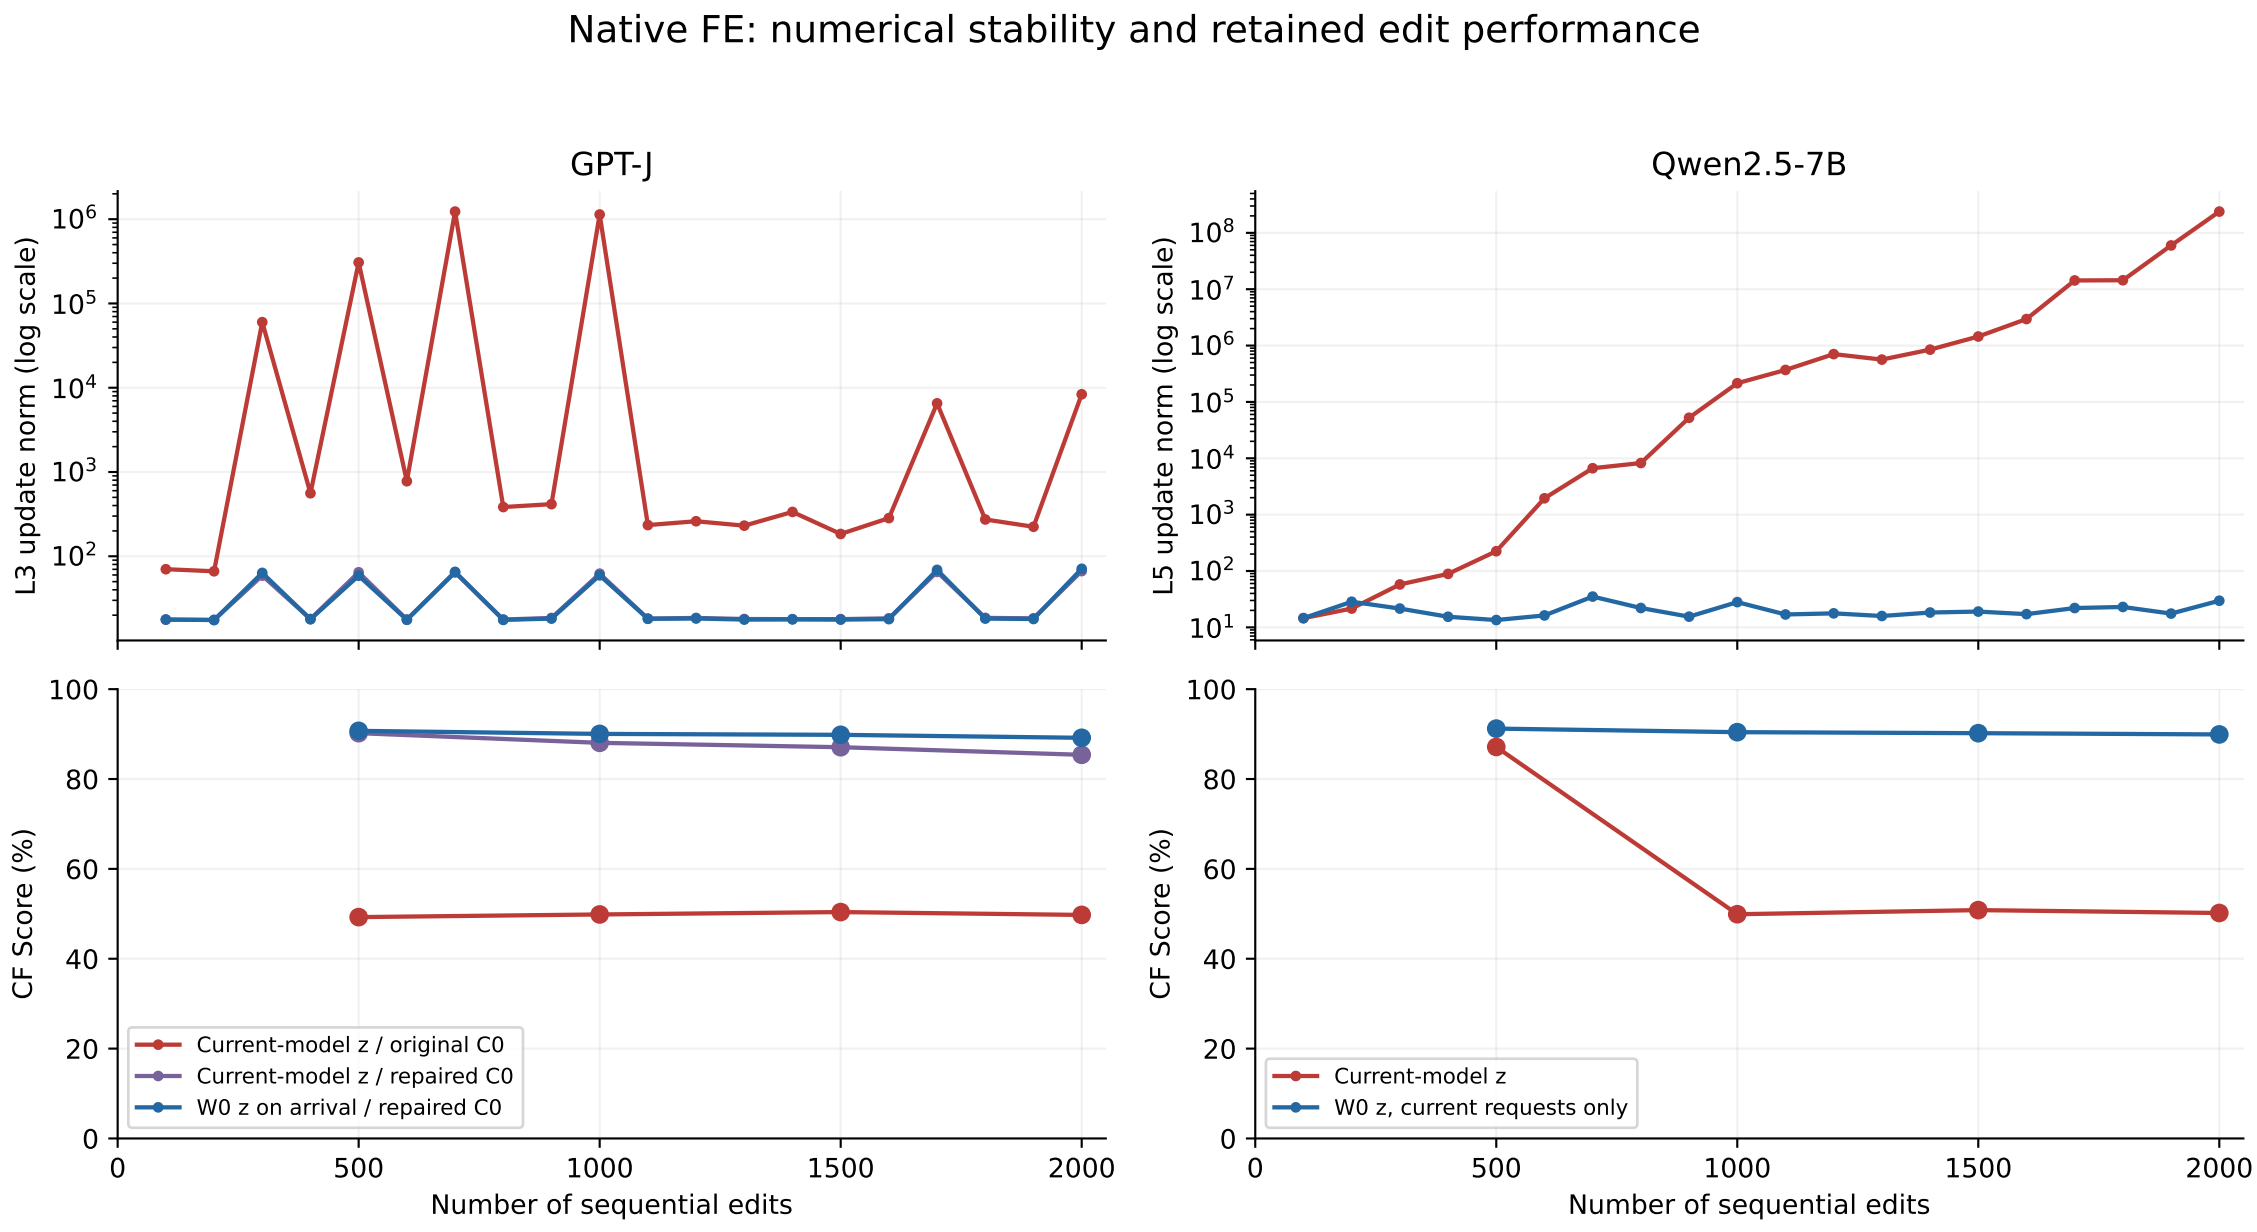
Native FE: numerical stability and retained edit performance
GPT-J
Qwen2.5-7B
L3 update norm (log scale)
L5 update norm (log scale)
CF Score (%)
Number of sequential edits
Current-model z / original C0
Current-model z / repaired C0
W0 z on arrival / repaired C0
Current-model z
W0 z, current requests only
Same ordered CF 2K stream; batch size 100. Score uses all seen requests at each checkpoint.
Single-stream comparison; W0 reference is a separately labeled protocol for both models.