



JLZ strength review at W10: 1,000 edits (BS100 × 10)

Same first 1,000 requests; success rates use the recorded R / P / N comparison metric.

(a) W10 comparison

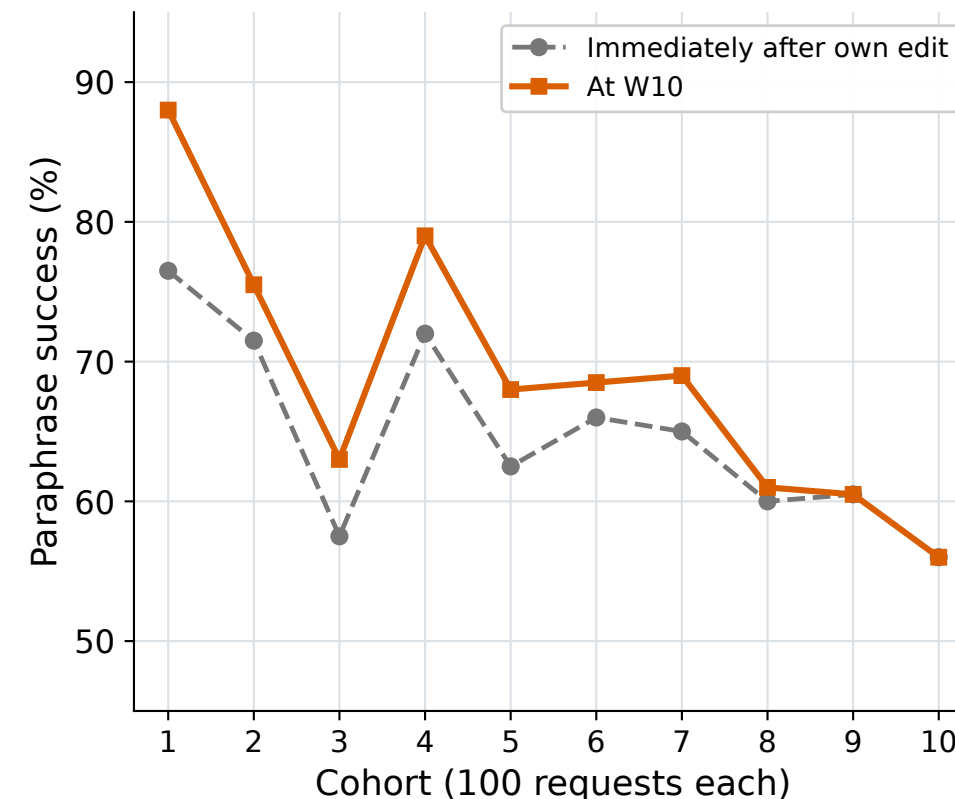


(b) JLZ-A: same cohort over time
200 paraphrases per cohort



Pooled: 64.15% → 67.30% (+3.15 pp)
Lost 26; gained 89 / 2,000 paraphrases

(c) JLZ-B: same cohort over time
200 paraphrases per cohort



Pooled: 64.75% → 68.85% (+4.10 pp)
Lost 13; gained 95 / 2,000 paraphrases

Birth = evaluation immediately after the cohort's own batch commit. W10 = the snapshot after batch 10.
Cohorts differ in difficulty; these curves compare each cohort with itself, not a fixed cohort across batch index.